

Capstone 2project Health Care (1)

May 16, 2023

[]:

```
[1]: #import the libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
import seaborn as sns
warnings.filterwarnings(action = "ignore", category = FutureWarning)
```

```
[2]: data= pd.read_csv("health care diabetes.csv")
data
```

```
[2]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	
..	
763	10	101	76	48	180	32.9	
764	2	122	70	27	0	36.8	
765	5	121	72	23	112	26.2	
766	1	126	60	0	0	30.1	
767	1	93	70	31	0	30.4	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0

```
766          0.349  47      1
767          0.315  23      0
```

```
[768 rows x 9 columns]
```

0.0.1 1. Perform descriptive analysis. It is very important to understand the variables and corresponding values. We need to think through - Can minimum value of below listed columns be zero (0)? On these columns, a value of zero does not make sense and thus indicates missing value. Glucose, BloodPressure, SkinThickness, Insuline, BMI. How will you treat these values?

```
[3]: data.describe()
```

```
[3]:      Pregnancies      Glucose  BloodPressure  SkinThickness      Insulin  \
count    768.000000    768.000000    768.000000    768.000000    768.000000
mean       3.845052    120.894531     69.105469     20.536458     79.799479
std        3.369578     31.972618     19.355807     15.952218    115.244002
min         0.000000     0.000000     0.000000     0.000000     0.000000
25%         1.000000     99.000000     62.000000     0.000000     0.000000
50%         3.000000    117.000000     72.000000     23.000000     30.500000
75%         6.000000    140.250000     80.000000     32.000000    127.250000
max        17.000000    199.000000    122.000000     99.000000    846.000000

      BMI  DiabetesPedigreeFunction      Age      Outcome
count    768.000000              768.000000    768.000000    768.000000
mean      31.992578              0.471876     33.240885     0.348958
std       7.884160              0.331329     11.760232     0.476951
min        0.000000              0.078000     21.000000     0.000000
25%       27.300000              0.243750     24.000000     0.000000
50%       32.000000              0.372500     29.000000     0.000000
75%       36.600000              0.626250     41.000000     1.000000
max       67.100000              2.420000     81.000000     1.000000
```

```
[4]: data.head(15)
```

```
[4]:      Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
0              6      148           72           35         0  33.6
1              1       85           66           29         0  26.6
2              8      183           64           0          0  23.3
3              1       89           66           23        94  28.1
4              0      137           40           35       168  43.1
5              5      116           74           0          0  25.6
6              3       78           50           32        88  31.0
7             10      115            0           0          0  35.3
8              2      197           70           45       543  30.5
9              8      125           96           0          0   0.0
```

10	4	110	92	0	0	37.6
11	10	168	74	0	0	38.0
12	10	139	80	0	0	27.1
13	1	189	60	23	846	30.1
14	5	166	72	19	175	25.8

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
5	0.201	30	0
6	0.248	26	1
7	0.134	29	0
8	0.158	53	1
9	0.232	54	1
10	0.191	30	0
11	0.537	34	1
12	1.441	57	0
13	0.398	59	1
14	0.587	51	1

0.0.2 we can see there are 0 in the columns of bloodpressure, bmi, insulin etc.. so this value doesnot make any sense. so replacing all the 0 value with median.

```
[5]: data['Glucose']=data['Glucose'].replace(0,data['Glucose'].median())
data['BloodPressure']=data['BloodPressure'].replace(0,data['BloodPressure'].
↳median())
data['SkinThickness']=data['SkinThickness'].replace(0,data['SkinThickness'].
↳median())
data['Insulin']=data['Insulin'].replace(0,data['Insulin'].median())
data['BMI']=data['BMI'].replace(0,data['BMI'].median())
data.describe()
```

```
[5]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin \
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	121.656250	72.386719	27.334635	94.652344
std	3.369578	30.438286	12.096642	9.229014	105.547598
min	0.000000	44.000000	24.000000	7.000000	14.000000
25%	1.000000	99.750000	64.000000	23.000000	30.500000
50%	3.000000	117.000000	72.000000	23.000000	31.250000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	32.450911	0.471876	33.240885	0.348958
std	6.875366	0.331329	11.760232	0.476951
min	18.200000	0.078000	21.000000	0.000000
25%	27.500000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000


0.0.3 2 Visually explore these variables using histograms. Treat the missing values accordingly.

```
[6]: data.isna().any()
```

```
[6]: Pregnancies      False
      Glucose          False
      BloodPressure    False
      SkinThickness    False
      Insulin          False
      BMI              False
      DiabetesPedigreeFunction  False
      Age              False
      Outcome          False
      dtype: bool
```


```
[7]: plt.hist("BloodPressure", data= data)
```

```
[7]: (array([ 3.,  2., 35., 118., 261., 214., 105., 18., 10., 2.]),
      array([ 24. , 33.8, 43.6, 53.4, 63.2, 73. , 82.8, 92.6, 102.4,
              112.2, 122. ]),
      <BarContainer object of 10 artists>)
```




```
[8]: plt.hist("Glucose", data= data)
```

```
[8]: (array([ 4., 19., 87., 149., 166., 125., 88., 54., 44., 32.]),  
      array([ 44. , 59.5, 75. , 90.5, 106. , 121.5, 137. , 152.5, 168. ,  
            183.5, 199. ]),  
      <BarContainer object of 10 artists>)
```




```
[9]: plt.hist("SkinThickness", data= data)
```

```
[9]: (array([ 59., 368., 181., 118.,  36.,   4.,   1.,   0.,   0.,   1.]),  
      array([ 7. , 16.2, 25.4, 34.6, 43.8, 53. , 62.2, 71.4, 80.6, 89.8, 99. ]),  
      <BarContainer object of 10 artists>)
```




```
[10]: plt.hist("Insulin", data= data)
```

```
[10]: (array([516., 143., 55., 29., 7., 10., 4., 1., 2., 1.]),  
      array([ 14. , 97.2, 180.4, 263.6, 346.8, 430. , 513.2, 596.4, 679.6,  
            762.8, 846. ]),  
      <BarContainer object of 10 artists>)
```



```
[11]: plt.hist("BMI", data= data)
```


```
[11]: (array([ 52., 161., 207., 193.,  91.,  48.,  10.,   4.,   1.,   1.]),  
      array([18.2 , 23.09, 27.98, 32.87, 37.76, 42.65, 47.54, 52.43, 57.32,  
            62.21, 67.1 ]),  
      <BarContainer object of 10 artists>)
```

```
[12]: fig, ax = plt.subplots(ncols= 3, figsize=(15,5))


sns.histplot(x= "Pregnancies", data= data, ax=ax[0])
sns.histplot(x= "Glucose", data= data, ax= ax[1])
sns.histplot(x= "BloodPressure", data= data, ax= ax[2])
```

```
[12]: <AxesSubplot:xlabel='BloodPressure', ylabel='Count'>
```



```
[13]: fig, ax= plt.subplots(ncols=2 , figsize=(15,5))
sns.histplot(x= "Insulin", data= data, ax= ax[0])
sns.histplot(x= "BMI", data= data, ax= ax[1])
```

```
[13]: <AxesSubplot:xlabel='BMI', ylabel='Count'>
```




0.0.4 3. There are integer and float data type variables in this dataset. Create a count (frequency) plot describing the data types and the count of variables.

```
[14]: data.dtypes
```

```
[14]: Pregnancies          int64
Glucose                  int64
BloodPressure            int64
SkinThickness            int64
Insulin                  float64
BMI                      float64
DiabetesPedigreeFunction float64
Age                     int64
Outcome                  int64
dtype: object
```


```
[15]: fig, ax= plt.subplots(ncols=3, figsize= (10,5))
sns.countplot(x= "Pregnancies", data= data, ax= ax[0])
sns.countplot(x= "Glucose", data= data, ax= ax[1])
sns.countplot(x= "BloodPressure", data= data, ax= ax[2])
plt.tight_layout()
```



0.0.5 4. Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of action.

```
[16]: sns.countplot(x= "Outcome", data= data)
```

```
[16]: <AxesSubplot:xlabel='Outcome', ylabel='count'>
```



```
[17]: print("value of \n", data["Outcome"].value_counts())
```

```
value of
0    500
1    268
Name: Outcome, dtype: int64
```

```
[18]: data["Outcome"].value_counts()/len(data)
```

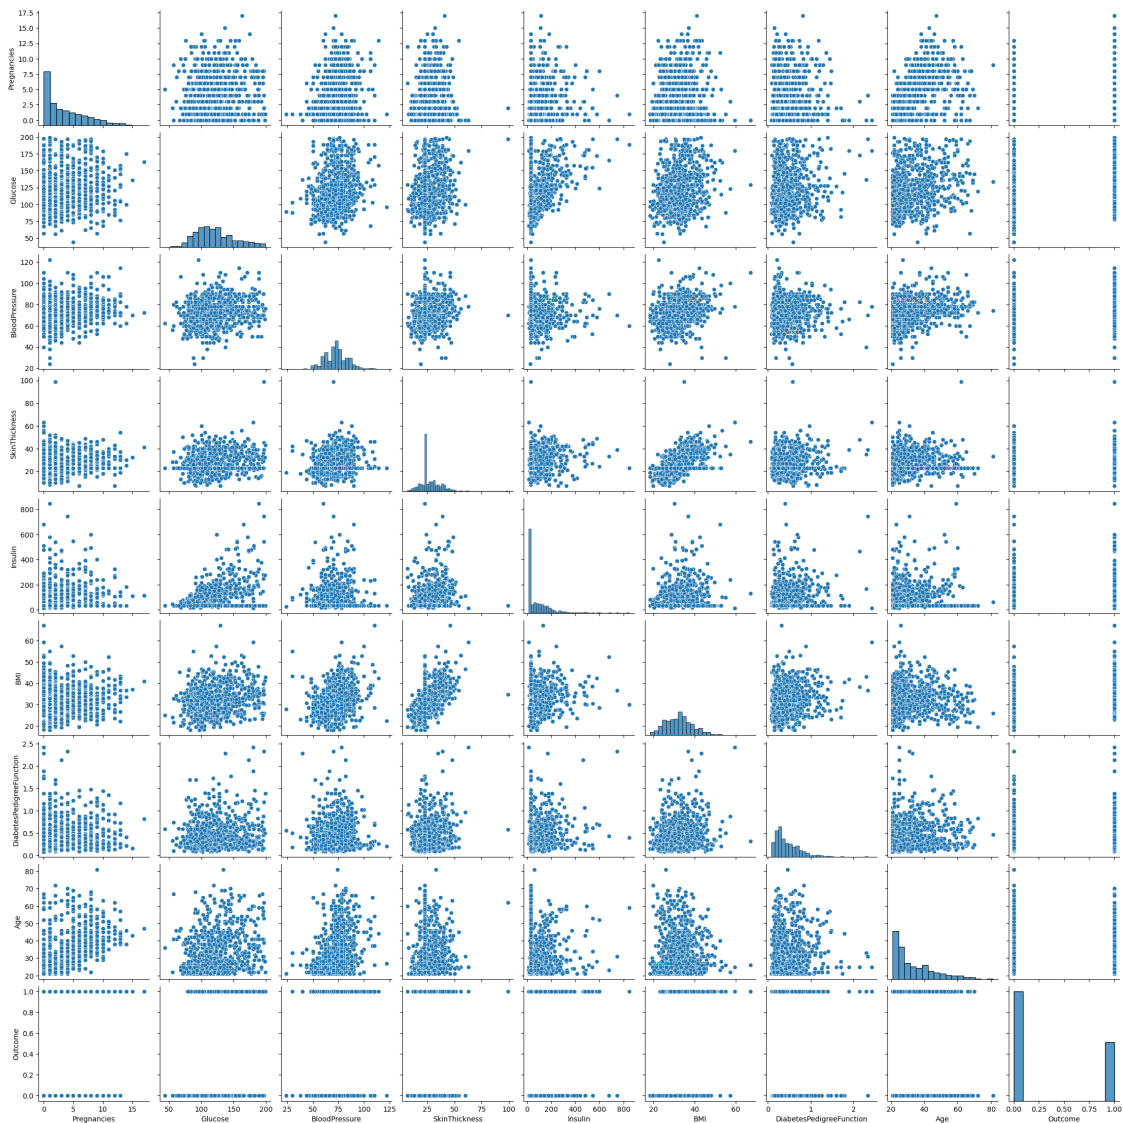
```
[18]: 0    0.651042
      1    0.348958
      Name: Outcome, dtype: float64
```

0.0.6 The data set is balanced.

0.0.7 5. Create scatter charts between the pair of variables to understand the relationships. Describe your findings.

```
[19]: sns.pairplot(data)
```

```
[19]: <seaborn.axisgrid.PairGrid at 0x2ae616e1550>
```



0.0.8 with this visualisation. we can see positive correlation between BMI and Skin Thickness ; and also between age and pregnancy,

```
[20]: data.corr()
```

```
[20]:
```


	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.128213	0.208615	0.032568	
Glucose	0.128213	1.000000	0.218937	0.172143	
BloodPressure	0.208615	0.218937	1.000000	0.147809	
SkinThickness	0.032568	0.172143	0.147809	1.000000	
Insulin	-0.055697	0.357573	-0.028721	0.238188	
BMI	0.021546	0.231400	0.281132	0.546951	
DiabetesPedigreeFunction	-0.033523	0.137327	-0.002378	0.142977	
Age	0.544341	0.266909	0.324915	0.054514	
Outcome	0.221898	0.492782	0.165723	0.189065	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.055697	0.021546	-0.033523	
Glucose	0.357573	0.231400	0.137327	
BloodPressure	-0.028721	0.281132	-0.002378	
SkinThickness	0.238188	0.546951	0.142977	
Insulin	1.000000	0.189022	0.178029	
BMI	0.189022	1.000000	0.153506	
DiabetesPedigreeFunction	0.178029	0.153506	1.000000	
Age	-0.015413	0.025744	0.033561	
Outcome	0.148457	0.312249	0.173844	

	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.266909	0.492782
BloodPressure	0.324915	0.165723
SkinThickness	0.054514	0.189065
Insulin	-0.015413	0.148457
BMI	0.025744	0.312249
DiabetesPedigreeFunction	0.033561	0.173844
Age	1.000000	0.238356
Outcome	0.238356	1.000000

```
[21]: plt.subplots(figsize=(15,9))
sns.heatmap(data.corr(), annot= True)
```

```
[21]: <AxesSubplot:>
```



1 Project Task: Week 2

1.1 Data Modeling:

1.1.1 1. Devise strategies for model building. It is important to decide the right validation framework. Express your thought process.

Apply an appropriate classification algorithm to build a model.

Compare various models with the results from KNN algorithm.

Create a classification report by analyzing sensitivity, specificity, AUC (ROC curve), etc.

Please be descriptive to explain what values of these parameter you have used.

1.1.2 The logistic regression will suit best as our dependent value is categorical data and independent variable are continuous.

```
[22]: #training and test data
X= data[["Pregnancies","Glucose", 'BloodPressure', 'SkinThickness', 'Insulin',
        'BMI', 'DiabetesPedigreeFunction', 'Age']]
y= data[["Outcome"]]

#Importing "train_test-split" function to test the model
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
        random_state=42)
```

```
[23]: print("shape of train data is", X_train.shape)
print("shape of test data is", X_test.shape)
```

```
shape of train data is (537, 8)
shape of test data is (231, 8)
```

```
[24]: #importing Logistic Regression
from sklearn.linear_model import LogisticRegression

lr = LogisticRegression()

#Fit the model in train and test data
lr.fit(X_train,y_train).score(X_train,y_train)
```

```
C:\Users\03man\anaconda3\lib\site-packages\sklearn\utils\validation.py:993:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
y = column_or_1d(y, warn=True)
C:\Users\03man\anaconda3\lib\site-
packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
[24]: 0.7877094972067039
```



```
[25]: #Now fitting the model in test set
prediction=lr.predict(X_test)
```

```
[26]: #Printing first 5 rows after fitting the model in test set
print (X_test.head())
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
668	6	98	58	33	190.0	34.0	
324	2	112	75	32	30.5	35.7	
624	2	108	64	23	30.5	30.8	
690	8	107	80	23	30.5	24.6	
473	7	136	90	23	30.5	29.9	

	DiabetesPedigreeFunction	Age
668	0.430	43
324	0.148	21
624	0.158	21
690	0.856	34
473	0.210	50

```
[27]: from sklearn import metrics
cm = metrics.confusion_matrix(y_test, prediction)
print('\n', "Confusion metrics is ", '\n', cm, '\n')
accuracy = metrics.accuracy_score(y_test, prediction)
print( '\n', "Accuracy score of logistic regression is :", accuracy, '\n')

print ( "classification score is", '\n', metrics.classification_report(y_test,
↪prediction))
```

```
Confusion metrics is
[[125  26]
 [ 31  49]]
```

```
Accuracy score of logistic regression is : 0.7532467532467533
```


```
classification score is
```

	precision	recall	f1-score	support
0	0.80	0.83	0.81	151
1	0.65	0.61	0.63	80
accuracy			0.75	231
macro avg	0.73	0.72	0.72	231
weighted avg	0.75	0.75	0.75	231

1.1.3 Logistic Regression gives 75% accuracy

```
[28]: from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve
lr_roc_auc = roc_auc_score(y_test, lr.predict(X_test))
fpr, tpr, thresholds = roc_curve(y_test, lr.predict_proba(X_test)[:,1])
plt.figure()
plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % lr_roc_auc)
plt.plot([0, 1], [0, 1], 'r--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic')
plt.legend(loc="lower right")
plt.savefig('Log_ROC')
print('AUC: %.3f' % lr_roc_auc)
plt.show()
```

AUC: 0.720



2 SVM Model

```
[29]: from sklearn.model_selection import train_test_split, GridSearchCV
      from sklearn.svm import SVC
      from sklearn.metrics import accuracy_score
```


```
[30]: model=SVC()
      model.fit(X_train, y_train)
      #creating the predictions on the test data i.e. X_test
      y_pred =model.predict(X_test)
      accuracy= accuracy_score(y_test, y_pred)
      print("Accuracy for the SVM Classifier: {:.3f}".format(accuracy)) #.3f is upto
      ↪3 places decimal
```

Accuracy for the SVM Classifier: 0.740

C:\Users\03man\anaconda3\lib\site-packages\sklearn\utils\validation.py:993:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
y = column_or_1d(y, warn=True)

```
[31]: from sklearn.metrics import roc_auc_score
      from sklearn.metrics import roc_curve
      lr_roc_auc = roc_auc_score(y_test, lr.predict(X_test))
      fpr, tpr, thresholds = roc_curve(y_test, lr.predict_proba(X_test)[:,:1])
      plt.figure()
      plt.plot(fpr, tpr, label='Logistic Regression (area = %0.2f)' % lr_roc_auc)
      plt.plot([0, 1], [0, 1], 'r--')
      plt.xlim([0.0, 1.0])
      plt.ylim([0.0, 1.05])
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('Receiver operating characteristic')
      plt.legend(loc="lower right")
      plt.savefig('Log_ROC')
      print('AUC: %.3f' % lr_roc_auc)
      plt.show()
```

AUC: 0.720



3 KNN Model

```
[32]: #feature Scaling
from sklearn.preprocessing import StandardScaler
st_X= StandardScaler()
X_train= st_X.fit_transform(X_train)
X_test= st_X.transform(X_test)
```

```
[33]: #Fitting K-NN classifier to the training set
from sklearn.neighbors import KNeighborsClassifier
classifier= KNeighborsClassifier(n_neighbors=5, metric='minkowski', p=2 )
classifier.fit(X_train, y_train)
```

```
C:\Users\03man\anaconda3\lib\site-
packages\sklearn\neighbors\_classification.py:198: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape
of y to (n_samples,), for example using ravel().
    return self._fit(X, y)
```

```
[33]: KNeighborsClassifier()
```

```
[34]: #Predicting the test set result
y_pred= classifier.predict(X_test)
```

```
[35]: #Creating the Confusion matrix
from sklearn.metrics import confusion_matrix , classification_report
cm= confusion_matrix(y_test, y_pred)
cm
```

```
[35]: array([[120,  31],
          [ 34,  46]], dtype=int64)
```

```
[36]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.78	0.79	0.79	151
1	0.60	0.57	0.59	80
accuracy			0.72	231
macro avg	0.69	0.68	0.69	231
weighted avg	0.72	0.72	0.72	231

```
[37]: accuracy= accuracy_score(y_test, y_pred)
print("Accuracy for KNN model is: {:.2f}".format(accuracy)) #.2f is upto 2
↳places decimal
```

Accuracy for KNN model is: 0.72

4 Random Forest

```
[38]: #fitting model
from sklearn.ensemble import RandomForestClassifier
clf= RandomForestClassifier(criterion="gini", #gini is entropy , y we use? we
↳see how much gini index reducing for this
                                max_depth= 7,
                                n_estimators= 200,
                                random_state= 5)

#the min. no. of trees in the forest should be atleast 40-50 , so n_estimator
↳=200 or nything more then 50 u can take
# the gini index heps us to understand the highest IG or the lowest entropy of
↳each feature. So essentially
# it will help us select the most imp variabllles in our model.
```

```
# max depth allow is 3, 5, 7)
```

```
[39]: #fitting the training data
      clf.fit(X_train, y_train)
```

```
C:\Users\03man\AppData\Local\Temp\ipykernel_30556\3181438177.py:2:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples,), for example using
ravel().
      clf.fit(X_train, y_train)
```

```
[39]: RandomForestClassifier(max_depth=7, n_estimators=200, random_state=5)
```

```
[40]: clf.feature_importances_
```

```
[40]: array([0.07197571, 0.32375323, 0.06285006, 0.06634026, 0.06758    ,
          0.16750927, 0.09165808, 0.14833338])
```

```
[41]: data.columns
```

```
[41]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
          'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
          dtype='object')
```

5 Glucose is very important as the value is max

```
[42]: y_pred= clf.predict(X_test)
```

```
[43]: from sklearn.metrics import confusion_matrix
      confusion_matrix(y_test, y_pred)
```

```
[43]: array([[125,  26],
          [ 31,  49]], dtype=int64)
```

```
[44]: #find accuracy score

      from sklearn.metrics import accuracy_score
      print("Accuracy of Random Forest is ", accuracy_score(y_test, y_pred))
```

```
Accuracy of Random Forest is  0.7532467532467533
```

6 decision Tree

```
[45]: #Fitting a decision tree classifier
from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier()
dtree.fit(X_train,y_train)
```

```
[45]: DecisionTreeClassifier()
```

```
[46]: #test the accuracy of the decision tree
predictions=dtree.predict(X_test)
from sklearn.metrics import classification_report,confusion_matrix
print(classification_report(y_test,predictions))
```

	precision	recall	f1-score	support
0	0.83	0.72	0.77	151
1	0.57	0.72	0.64	80
accuracy			0.72	231
macro avg	0.70	0.72	0.70	231
weighted avg	0.74	0.72	0.72	231

```
[47]: from sklearn.metrics import accuracy_score
print("Accuracy of Decision Tree is ", accuracy_score(y_test, predictions))
```

Accuracy of Decision Tree is 0.7186147186147186